

# Scalable Stream Processing and Map-Reduce

Neel Sundaresan, Evan Chiu, Gyanit Singh

eBay Research Labs



# About eBay Research Labs

- Who we are
  - eBay Research Labs was formed in July of 2005. The group's charter is to conduct forward looking research and deliver innovative solutions to business and product
- What we do
  - Search & IR
  - Machine learning
  - Analytics and Optimization
  - Reputation, Trust and Safety
  - Distributed and Grid Computing
  - Social and Incentive Networks
  - Large Scale Visualization
  - Scalable Martix and Graph Computing
  - ...
- Basically we “Dig” into data!

# The Land of Large Numbers

eBay users trade about \$1,400 worth of goods on the site every second.

On an average day on eBay...

shoe

sells


ever

y 10


minu

tes

# Listings



# Registered Users



# League of Long Tail



# Of Needles, Haystacks, and Magnets...

- Transaction Data and session data
  - Terabytes per day
- Data mining researcher's source of truth
- Nature of session data
  - Sessionized streams
  - Semi-structured
  - Constantly changing schema
- Examples:
  - View item click through rate for search algorithms.
  - Browsing pattern of users performing searches in different categories.
    - For e.g. computer vs clothing.
  - Purchase rate for various recommendation algorithms – A/B experimentation

# The Log Challenge

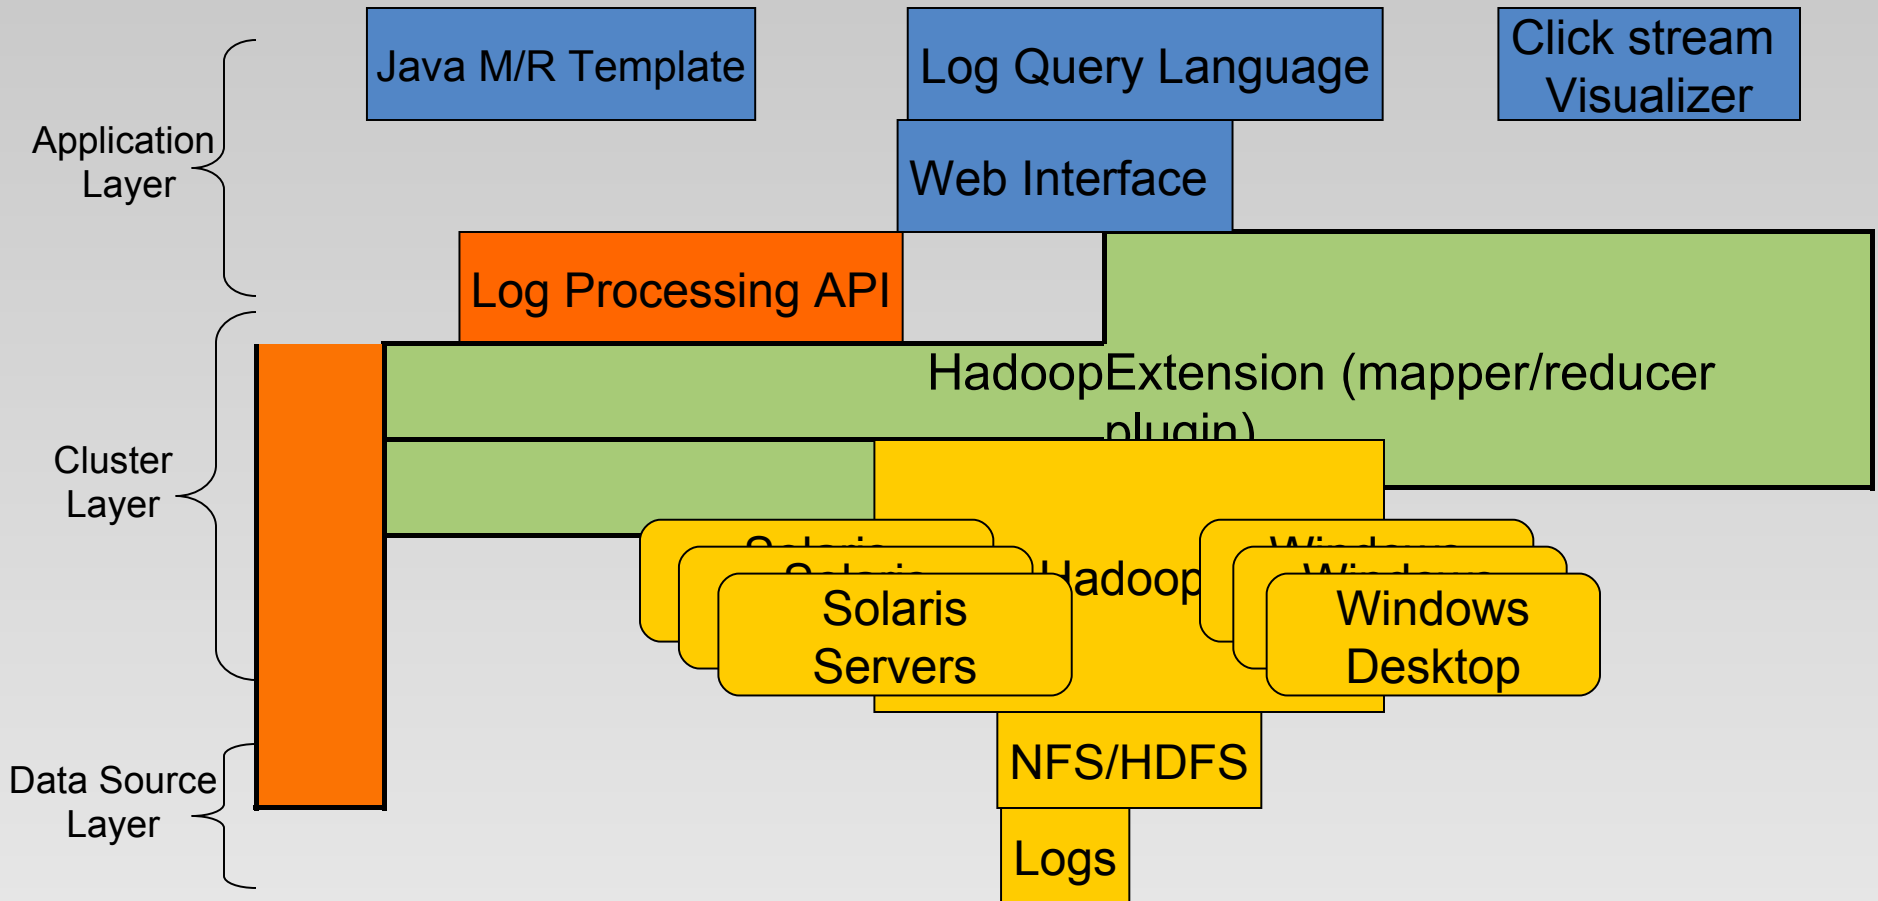
- The amount of the data is huge.
  - TB+ / day
  - Need to perform analysis on weeks or even years worth of data.
- Analysis takes time.
  - Making it distributed will help.
- Text streams are not indexed.
  - Difficult to query
  - Fields often change
- Difficult to perform sessionized analysis
  - E.g.: Study the session paths of session in which a given search algorithm is used.
  - E.g.: Differences in UK user and US users.
- Large Number of session paths (in millions)
  - Visualization is difficult
- Scaling up to potential large user base




# What we need is..

- An analytic tool:
  - Easy to perform session analysis.
  - Large scale stream processing.
  - Quick turnaround on analysis.
  - Effective query language.
  - Visual analytics that caters to intuition and provides extensive analytics.
  - Highly customizable processing.
  - Provide interfaces at different levels.

# Architecture - Mobius





# User session information

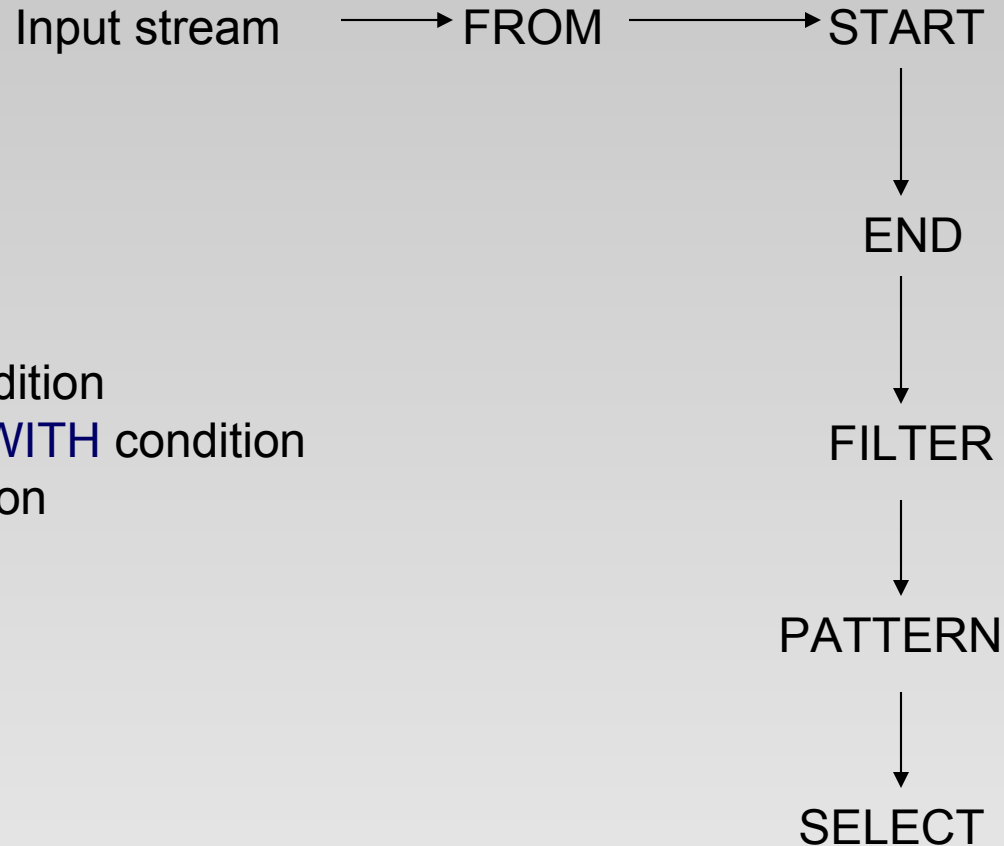
- Questions
  - What percentage of searches done receive clicks?
  - Out of those clicked results, how many are abandoned?
  - How many viewed results are followed to bid?
- Data
  - Session activity grouped together as a stream.
  - A session is a bag of events.
  - Each event is a tuple with various fields.
- Process:
  - Extract session with searches.
  - Compute view through rate.

# What do we need in a Stream Query Language

- Detect patterns from the stream over a window
  - window
    - time-based, count-based, event/trigger-based
    - Sliding vs Landmark windows
  - Structure
    - Repetition (\*, +), Sequence (//), condition
  - Naming patterns
- Integrate Relational and Pattern operators
- “Inner” Queries

# Mobius Query Language (MQL) - Structure of a Query

SELECT \*  
FROM source  
WHERE where\_condition  
PATTERN pattern WITH condition  
START start\_condition  
END end\_condition



# Simple Example

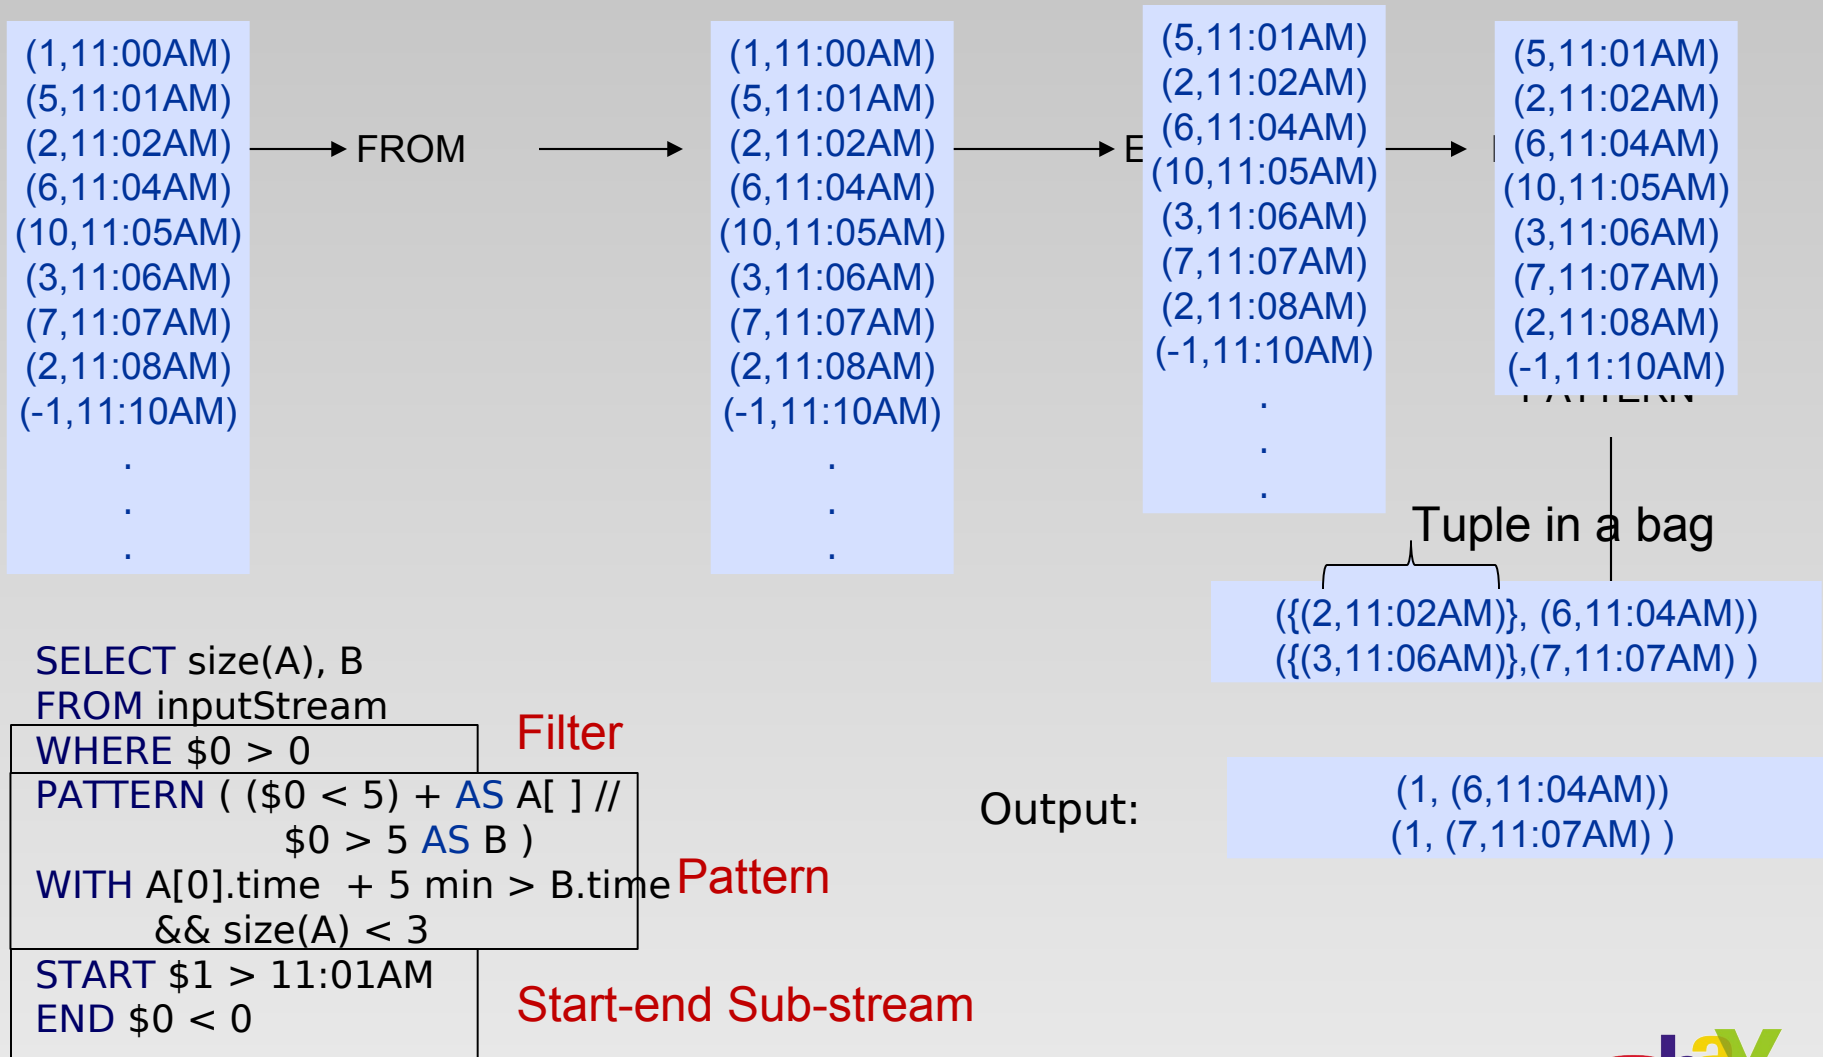
- Input: Stream of events.
- Event is modeled as tuples.
- Stream is modeled as ordered bag.

inputStream

(1,11:00AM)  
(5,11:01AM)  
(2,11:02AM)  
(6,11:04AM)  
(10,11:05AM)  
(3,11:06AM)  
(7,11:07AM)  
(2,11:08AM)  
(-1,11:10AM)

·  
·  
·

# Simple Example Contd...





# Other Systems

- PIG and Hive
  - Patterns are not available.
  - Other stream processing operator also not available. E.g. Start, End.
- CEP based Stream Processing Languages. (STREAM, Streambase, Cayuga)
  - Have flat data model.
    - Can only store few features of patterns.
  - Degree of parallelism is restricted to 1 due to inability to represent substreams.
    - In some cases splitting is done but that splitting operator has 1 degree of parallelism.
- Active Databases
  - ECA (events-conditions-actions) and triggers

# Pattern Query

- Problem: For each user identify a set of “search” followed by “view” (click-thru) events
- Input Stream : (uid,session(name,time, itemID,...))\*
- Output Stream : (uid,views(S,V))\*


- DATASET clicks = SELECT uid, {

```
SELECT S, V
FROM session
PATTERN (name == “search” AS S //
        ** //
        name == “view” AS V)
WITH V.itemID belongsto S.impressions
} AS views
FROM logs
```

```
WHERE size(session) > 1;
```

Pattern Query

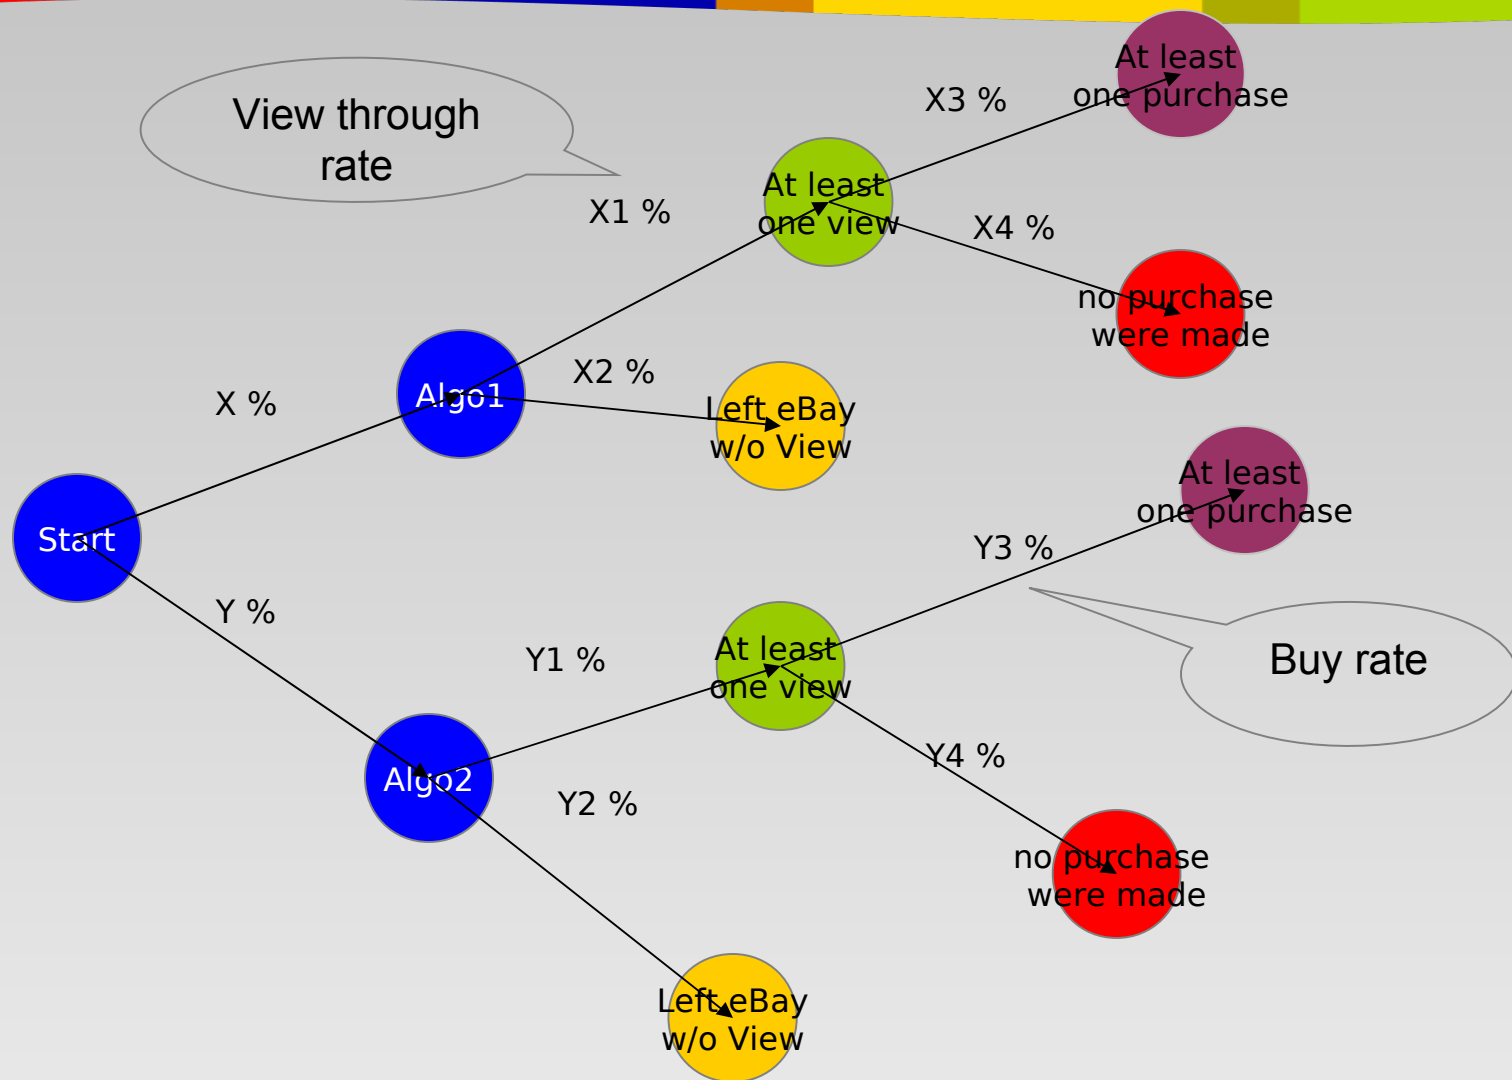
Pattern defined in the pattern query is



# Recommender Systems

- Products are recommended to users on various pages.
  - How many clicks does the recommendation gets?
  - Do those clicks result in purchase?
  - Performance of different recommendation algorithms?
- Data
  - User session containing events.
  - Events are of different types
- Process
  - Extract session with recommendations.
  - Group them by algorithm used.
  - Calculate the view and purchase through rate.

# Sessionized Analysis 1




# Step1 : Extract all sessions with “view”s


- Input Stream: merch\_logs is (uid, sessionid, (name,time,itemID..)\*)\*
- DATASET** merchView = **SELECT** uid, sessionid, {

```

SELECT S.algorithm AS algo, V.itemID AS itemID
FROM session
PATTERN ( name == “recommend” AS S//
           ** AS B[ ] //
           name == “view” AS V)
WITH (S.itemID == V.itemID) AND
       (B[i].itemid != S.itemID)
AS viewedRecos
FROM merch_logs ;
    
```



- The PATTERN part defines pattern




- Schema of merchView is (uid, sessionid, (algo,itemID)\*)\*

Bag named viewedRecos

## Step 2: Extract all instances of recommendations made...


- **DATASET** merchShown = **SELECT** uid, sessionid, flatten(  
    **SELECT** algorithm AS algo  
    **FROM** session  
    **WHERE** name == "recommend"  
    })  
    **FROM** merch\_logs;  
(uid, sessionid, (name,time,itemID...)\* => (uid,sessionId, algo)\*



# Step 3: Produce a flattened view


- DATASET unnestedmerchView = SELECT uid, sessionid, flatten(viewedRecos)  
FROM merchView;

(uid, sessionid, (algo,itemID)\*)\* => (uid, sessionid, algo, itemId)\*



# Step 4: Group the data by algorithm type and compute counts

- **DATASET** merchData = **SELECT** groupid **AS** algo, size(unnestedmerchView) **AS** clicks, size(merchShown) **AS** impressions  
**FROM** unnestedmerchView, merchShown  
**GROUP** unnestedmerchView **BY** algo  
**ALSO** merchShown **BY** algo

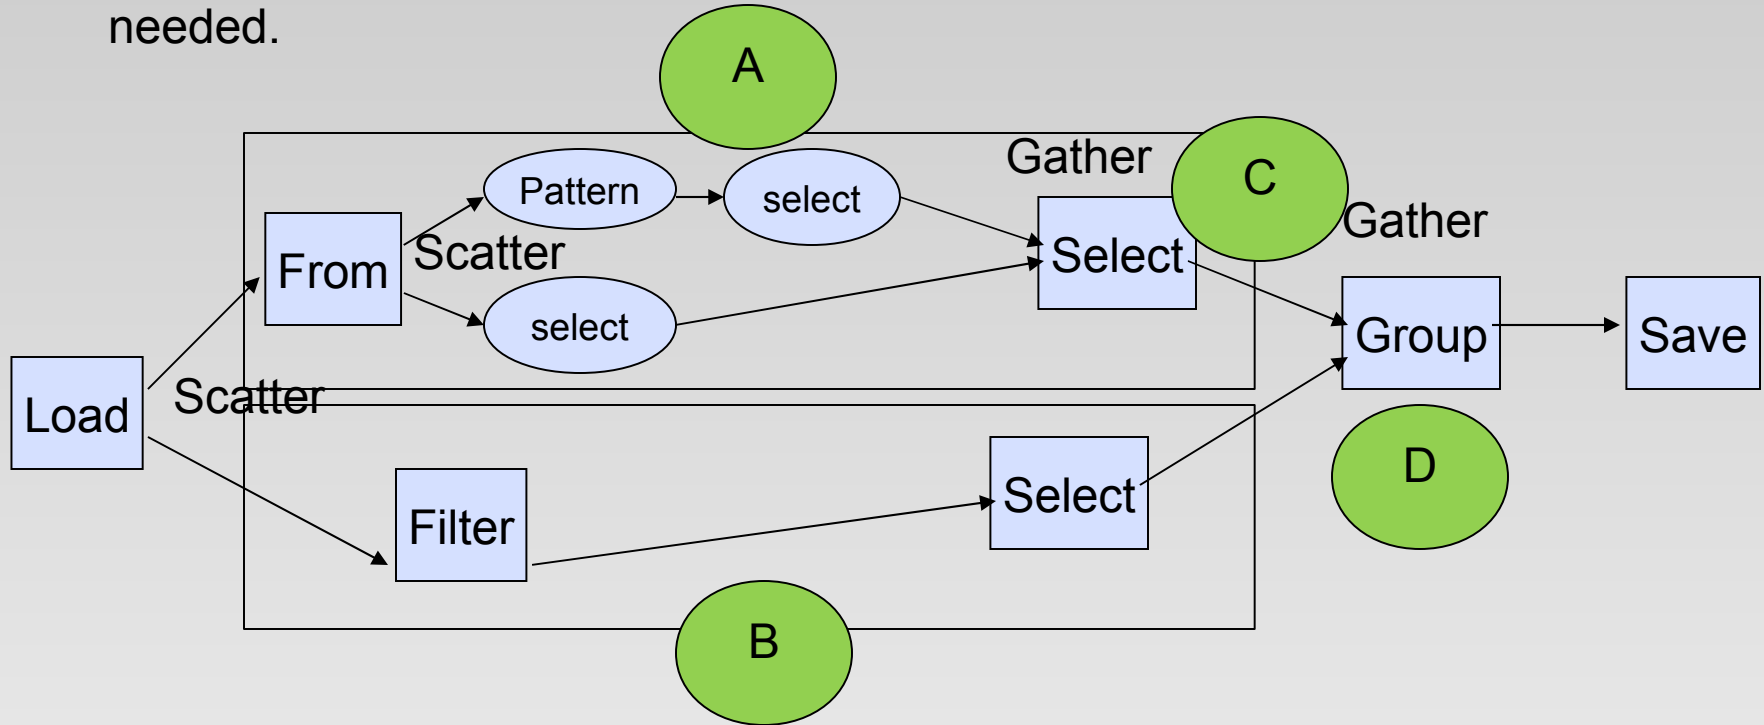


(uid, sessionid, algo, itemId)\* , (uid,sessionId, algo)\* => (algo, #clicks, #impression)\* -- stream/bag of length #algos




# Parallel Implementation

- MQL compiling engine compiles queries in to a DAG. (similar to PIG)
- Then the DAG is compiled into one or more map-reduce jobs.
- When ever a grouping, sorting, union operator is seen a reduce phase is needed.




# Understanding System Health

- Data
  - System logs containing many beacon streams.
  - Each machine beacon stream comprising of single beacon message.
  - Message contains various state data.
- Problem
  - Contiguous time period load is more than 60%.
  - Time period only interesting if it is more than delta mins.



- Input Stream: Schema of systemlogs is (systemname, beacons(load, time, ...))
- Output Stream: (systemname, (load,time)\*)\*
- DATASET system = 

```
SELECT systemname, {
    SELECT LE AS load
    FROM beacons
    PATTERN ( load < 60 AS SEVENT //
            (load > 60)+ AS LE[ ] //
            load < 60 AS EEVENT)
    WITH LE[size(LE)].time – LE[0].time > 10mins
    } AS LoadTimes
FROM systemlogs;
```
- Pattern defined in pattern query is



# Ongoing and Future Work

- Optimization mechanism.
  - Filter cannot be pushed ahead of user defined functions and pattern queries.
  - Inferring projection is also limited.
- Compilation to Map-Reduce Jobs
  - Inferring the best strategy to split work between map-reduce in case of multiple queries.
- Degree of parallelization when stream is not split by the users into substreams.
- Near real-time stream processing engine.
- Reporting mechanism.